



# Introduction to Computer Vision

## Lecture 10 - 3D Vision III and Temporal Analysis I

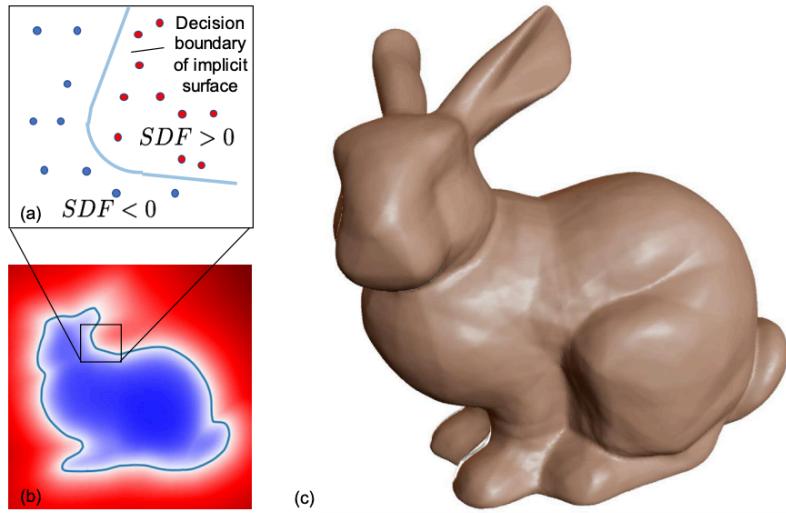
Prof. He Wang

# Logistics

- Assignment 3:
  - released on 4/26 evening
  - due on 5/10 11:59PM (Saturday)
- If 1 day (0 - 24 hours) past the deadline, 15% off
- If 2 day (24 - 48 hours) past the deadline, 30% off
- Zero credit if more than 2 days.

# Implicit Field

# Implicit Shape

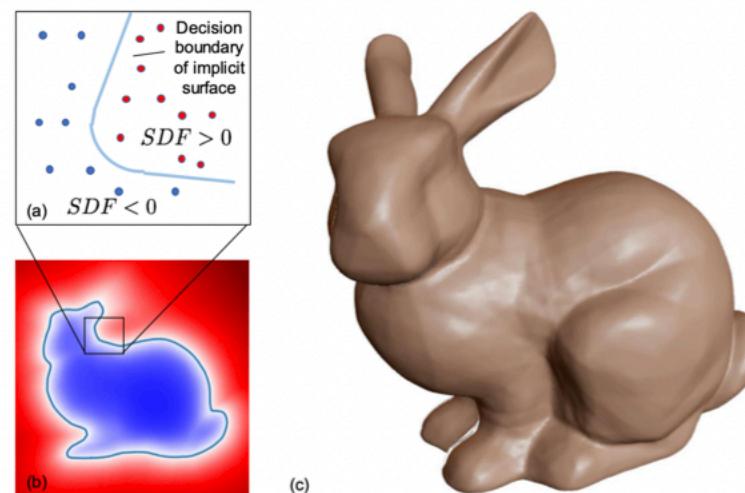


SDF

- Both an implicit geometry and surface representation
- Can convert into mesh
- Signed distance function, unsigned distance function, occupancy network

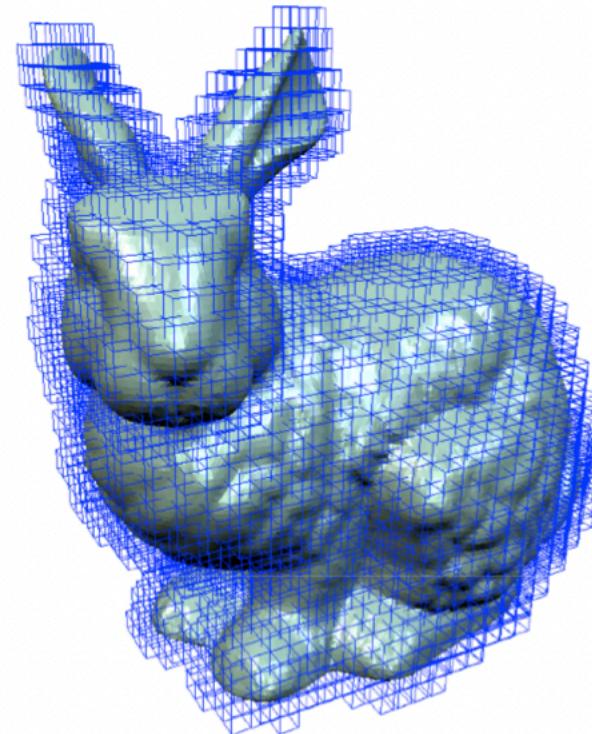
# Signed Distance Function (SDF)

- Interior:  $F(x, y, z) < 0$
- Exterior:  $F(x, y, z) > 0$
- Surface:  $F(x, y, z) = 0$  (zero set, zero iso-surface)
- Example implementation:
  - SDF:  $F(x, y, z) = \text{distance to the surface}$



# How to Extract Zero Iso-Surface

Input: a signed distance field  
(Implicitly assumed knowing the  
inside/outside of the shape, often  
needs to be estimated with  
**normal information**)



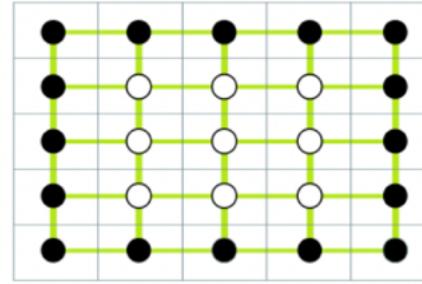
# Classical Solution: 2D Marching Square

1	1	1	1	1	1
1	2	3	2	1	
1	3	3	3	1	
1	2	3	2	1	
1	1	1	1	1	1

Threshold  
with iso-value  

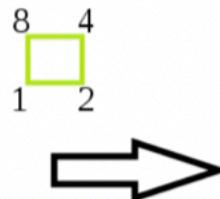

0	0	0	0	0	0
0	1	1	1	1	0
0	1	1	1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

Binary image  
to cells  

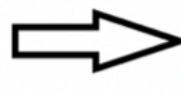



# Classical Solution: 2D Marching Square

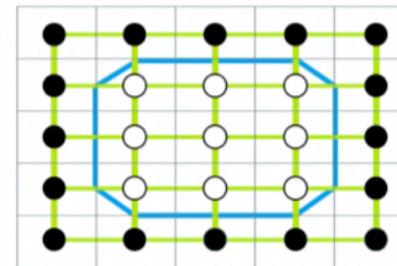
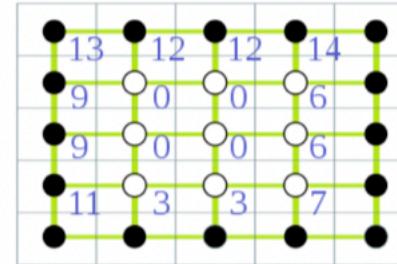
Give every cell a number based on which corners are true/false



Look up the contour lines in the database and put them in the cells

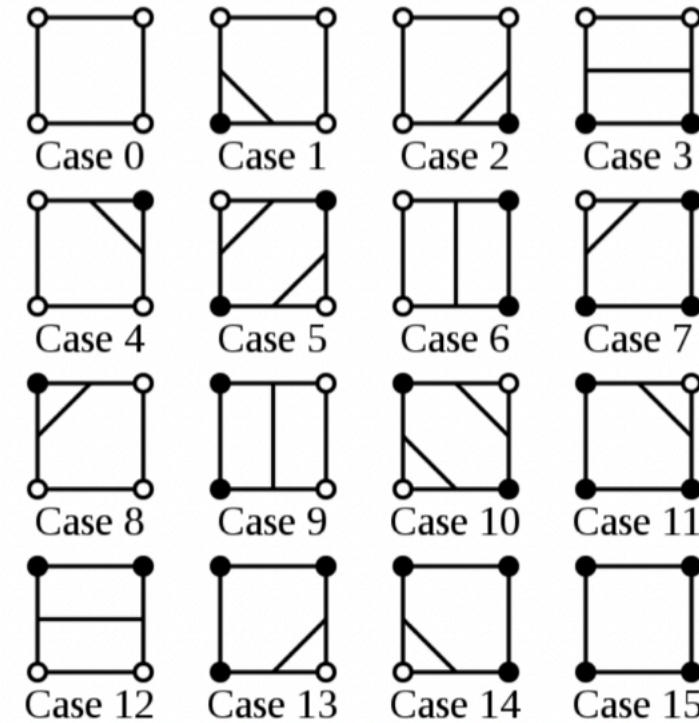


Look at the original values and use linear interpolation to determine a more accurate position of all the line end-points



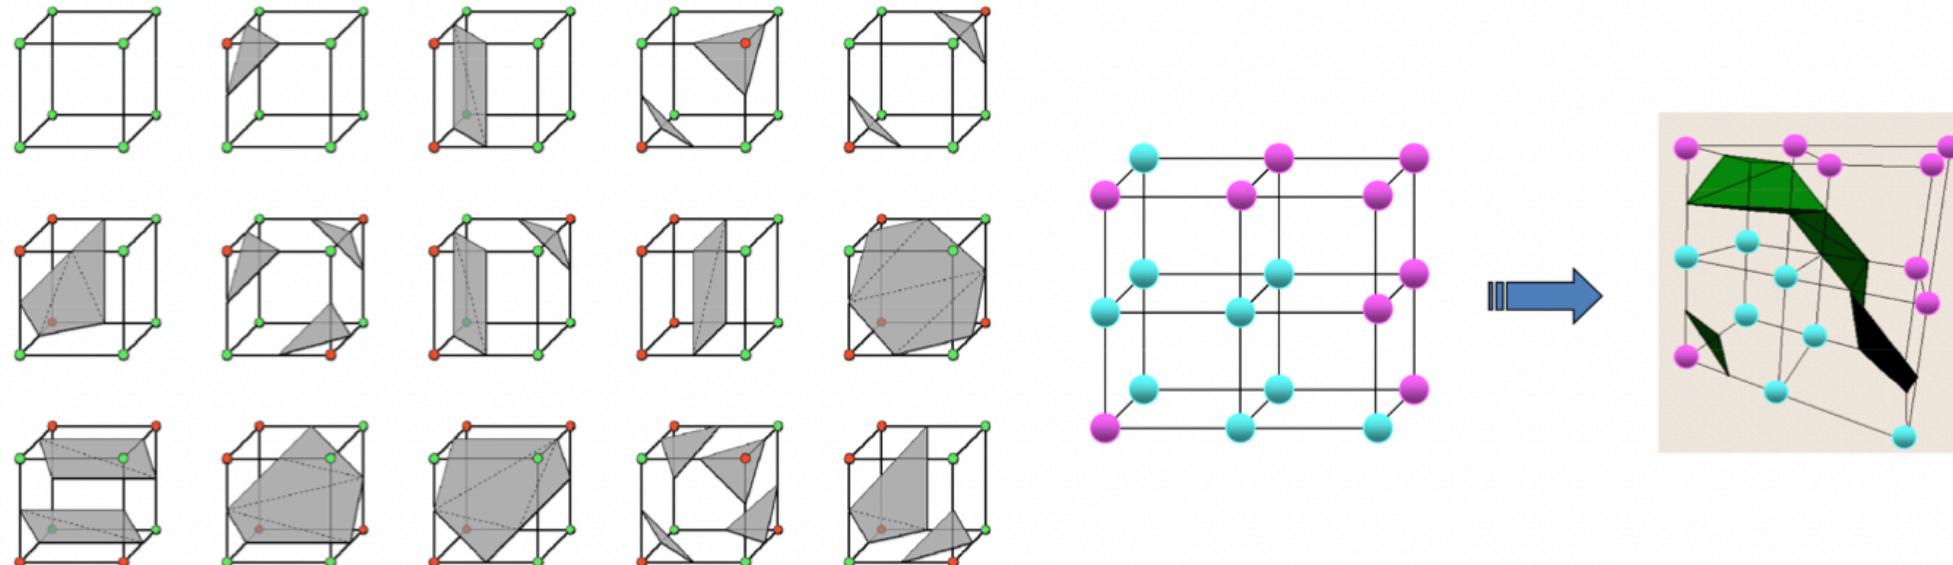
1	1	1	1	1
1	2	3	2	1
1	3	3	3	1
1	2	3	2	1
1	1	1	1	1

Look-up table contour lines



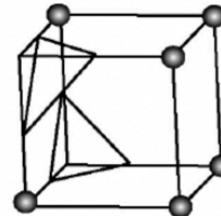
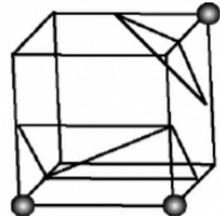
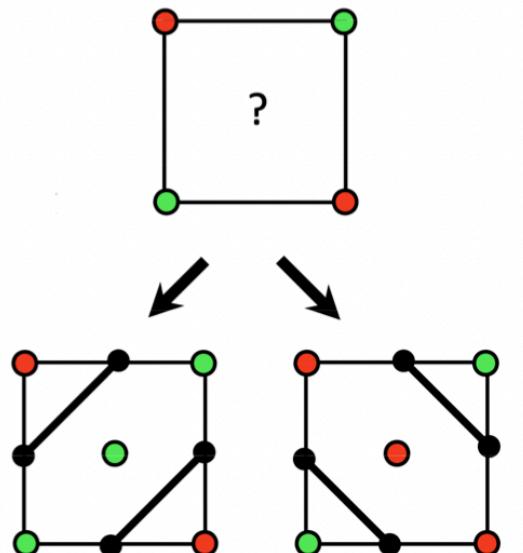
# Classical Solution: 3D Marching Cube

- $2^8 = 256$  cases
- The first published version exploits rotation and inversion, and only considers 15 unique cases:

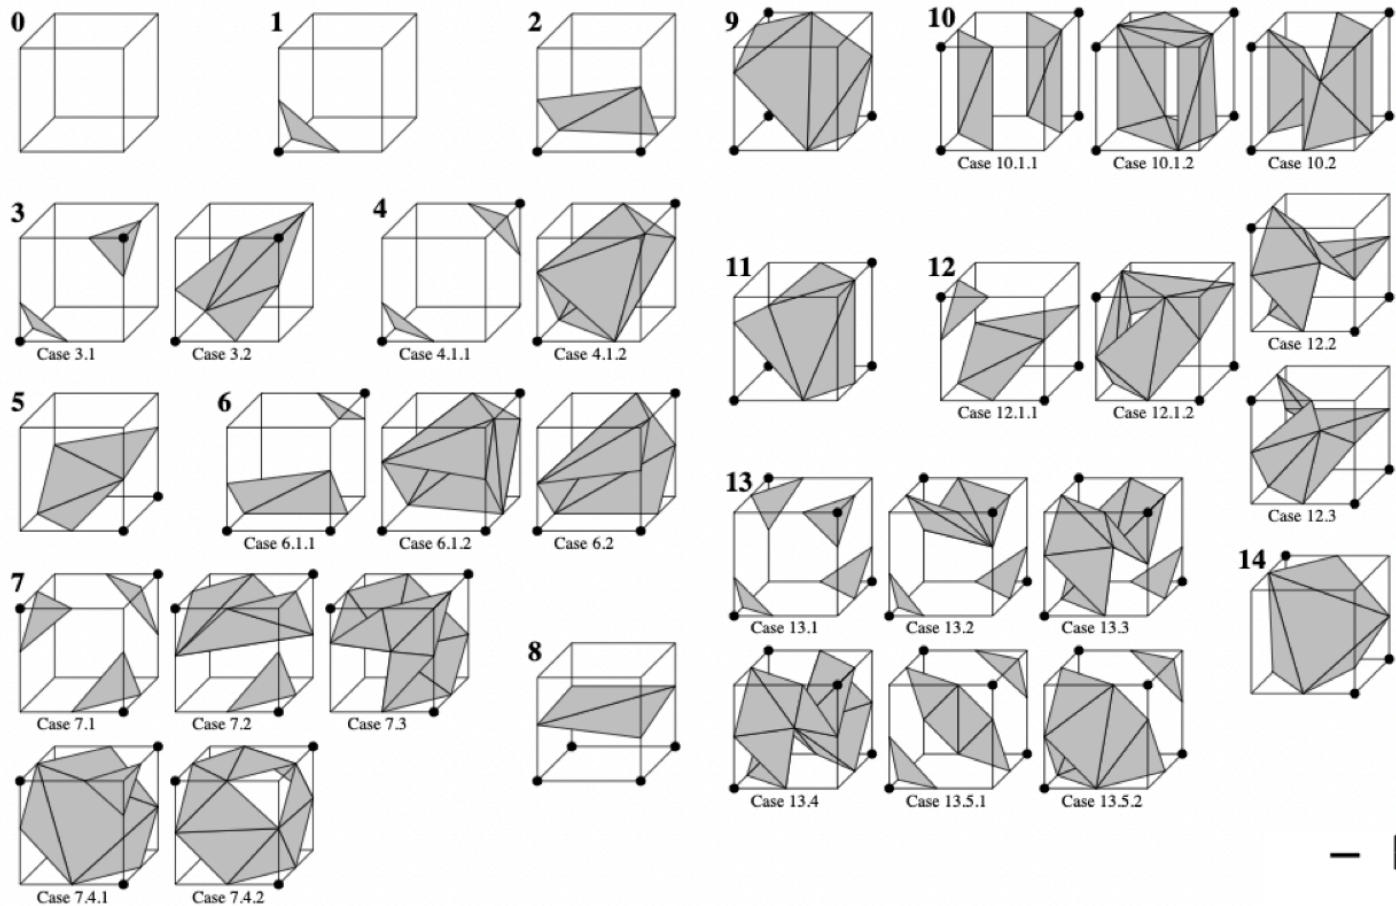


# Ambiguity

- Ambiguity leads to holes:



# Solution to Ambiguity



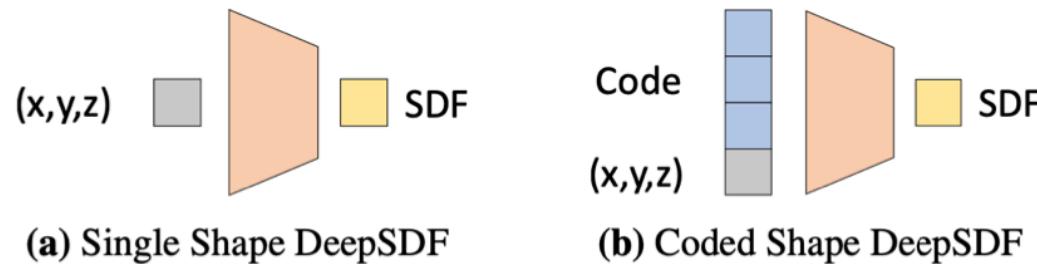
a

(a)

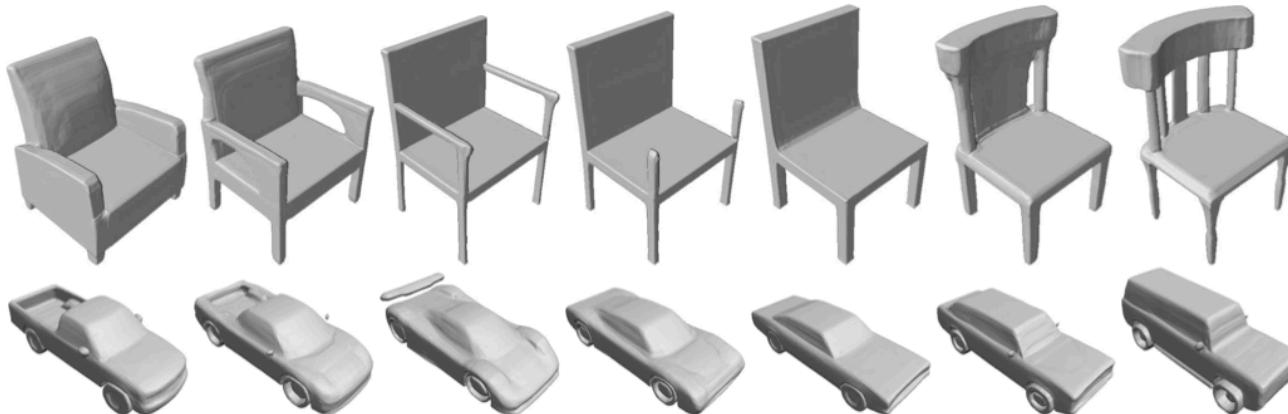
(b)

- Rotation only
- Always separate same “color”
- → Ambiguous faces triangulated consistently

# Deep SDF



- (a) use the network to overfit a single shape
- (b) use a latent code to represent a shape, so that the network can be used for multiple shapes



# Neural Radiance Field

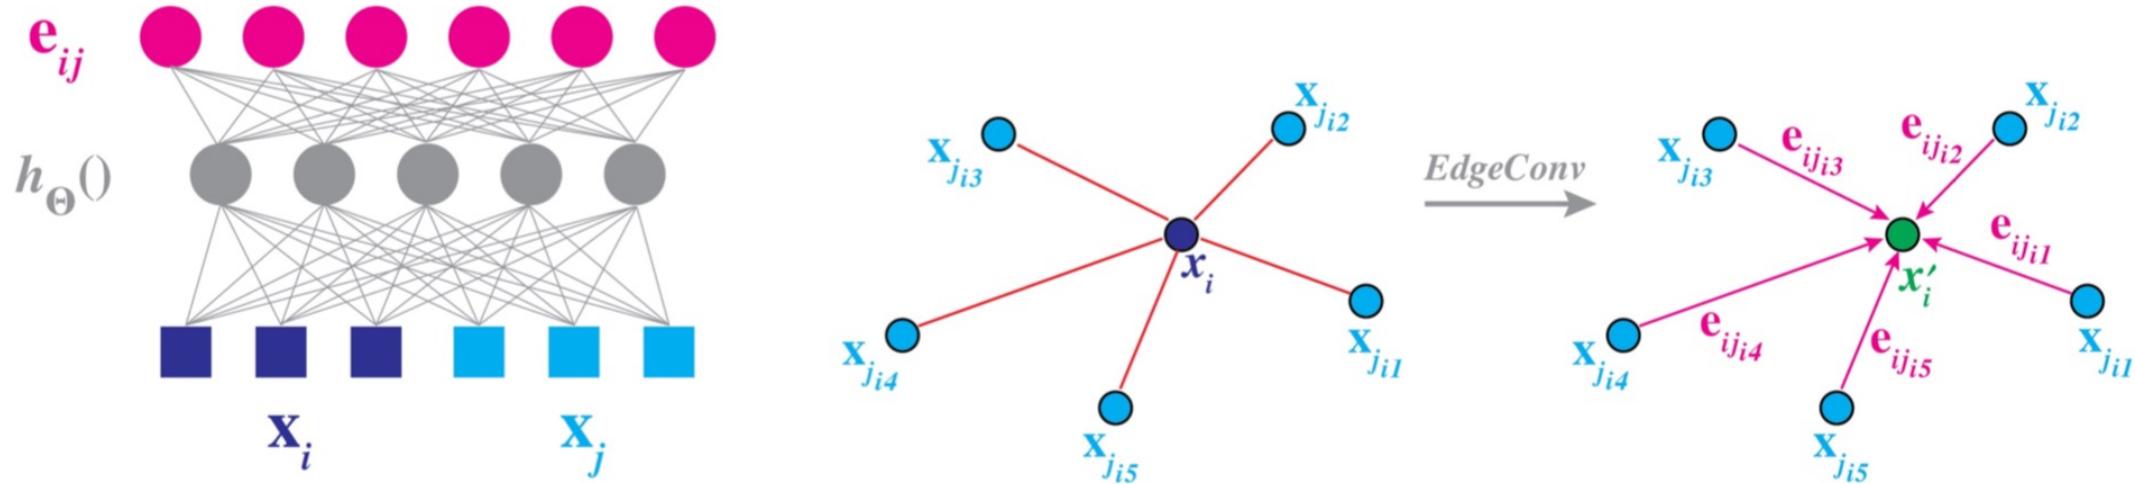
Training network to reproduce all input views of the scene



## Neural Radiance Field (NeRF)

A brief intro to NeRF: <https://www.youtube.com/watch?v=JuH79E8rdKc>

# Convolution on Mesh/Graph



**Message passing:** The output of EdgeConv at the  $i$ -th vertex is thus given by

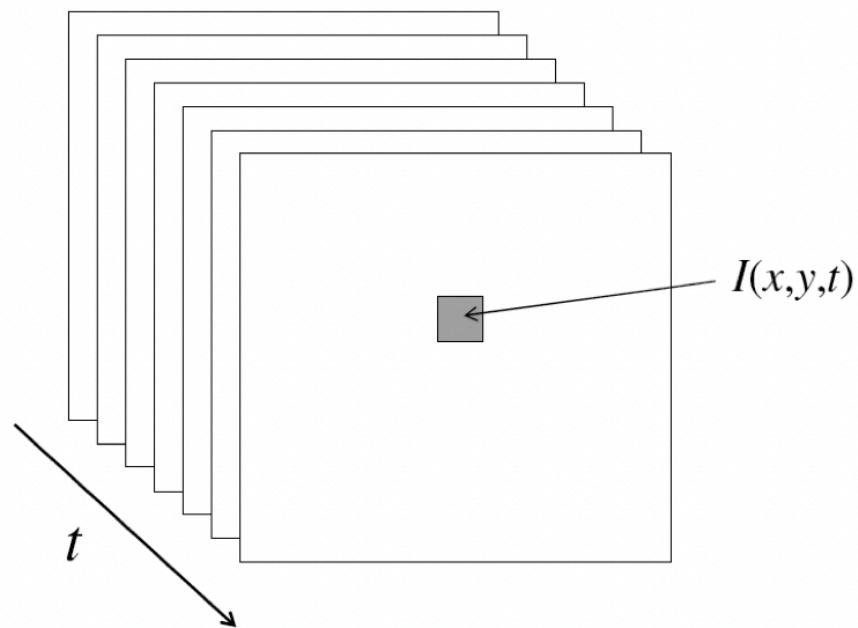
$$\mathbf{x}'_i = \square_{j:(i,j) \in \mathcal{E}} h_\Theta(\mathbf{x}_i, \mathbf{x}_j). \quad (1)$$

# Motion and Optical Flow

Some slides are borrowed from Stanford CS131.

# From Single Image to Video

- A video is a sequence of frames captured over time
- Now our image data is a function of space ( $x, y$ ) and time ( $t$ )



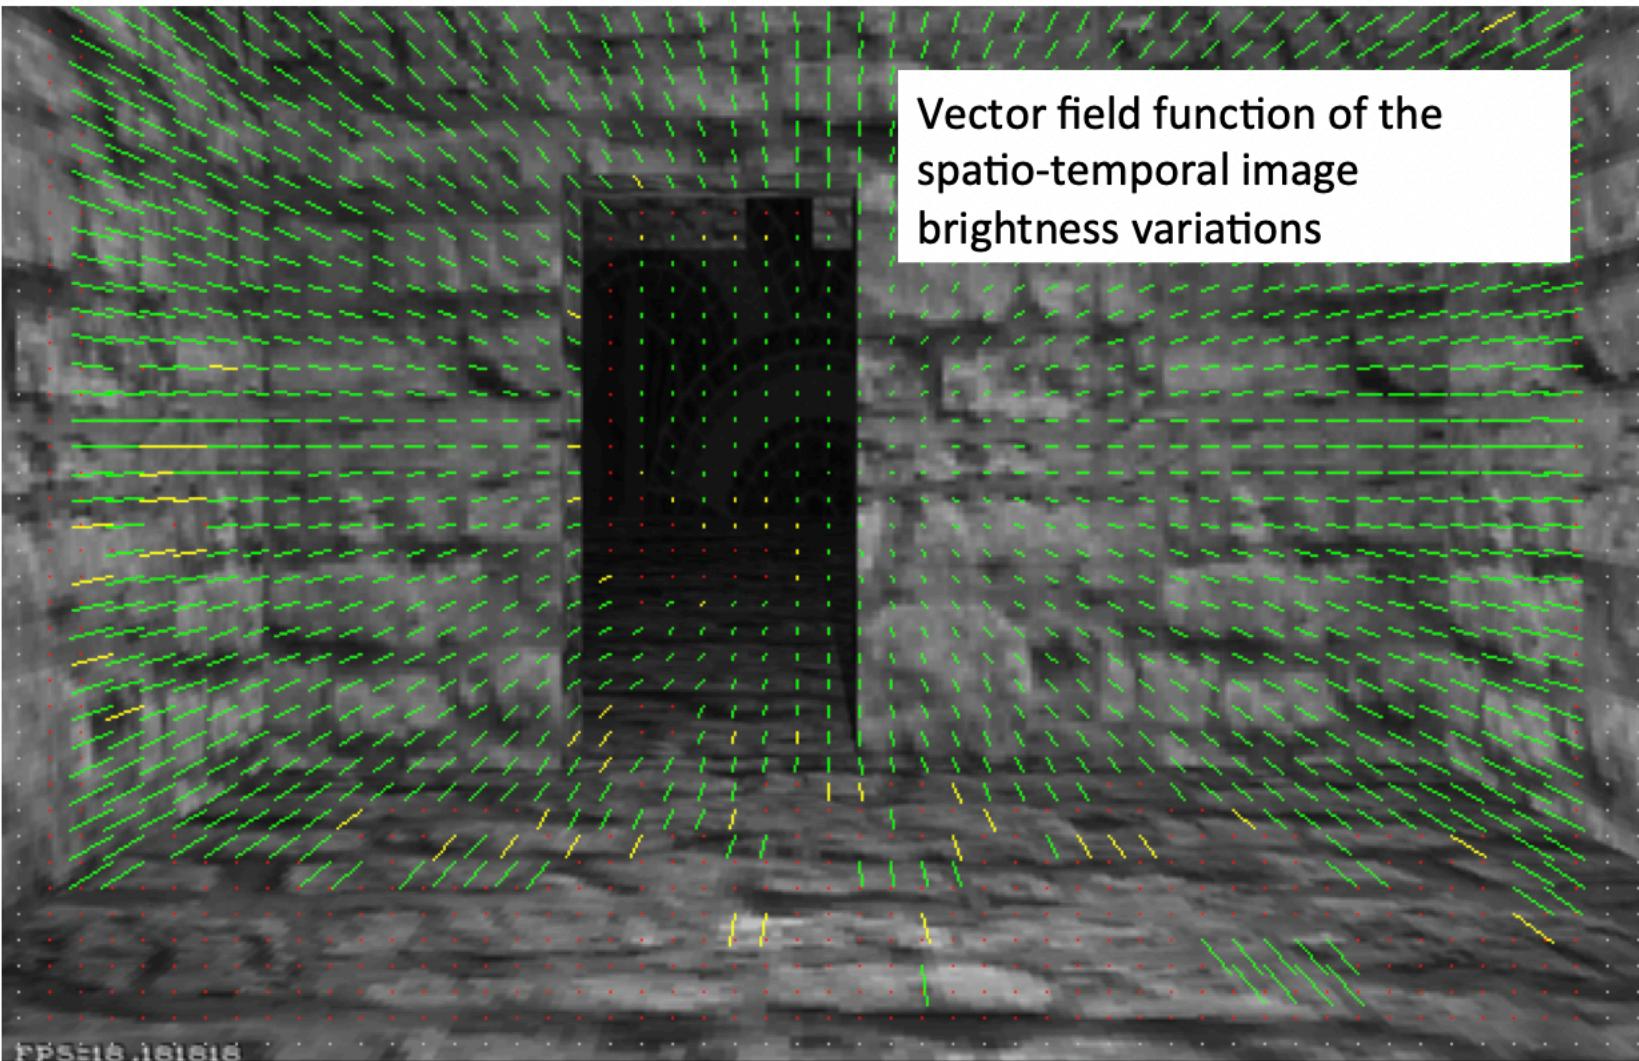
**Today let's focus on motions between two consecutive frames!**

# Optical Flow

- Definition: optical flow is the *apparent* motion of brightness patterns in the image
- Note: apparent motion can be caused by lighting changes without any actual motion
  - Think of a uniform rotating sphere under fixed lighting vs. a stationary sphere under moving illumination

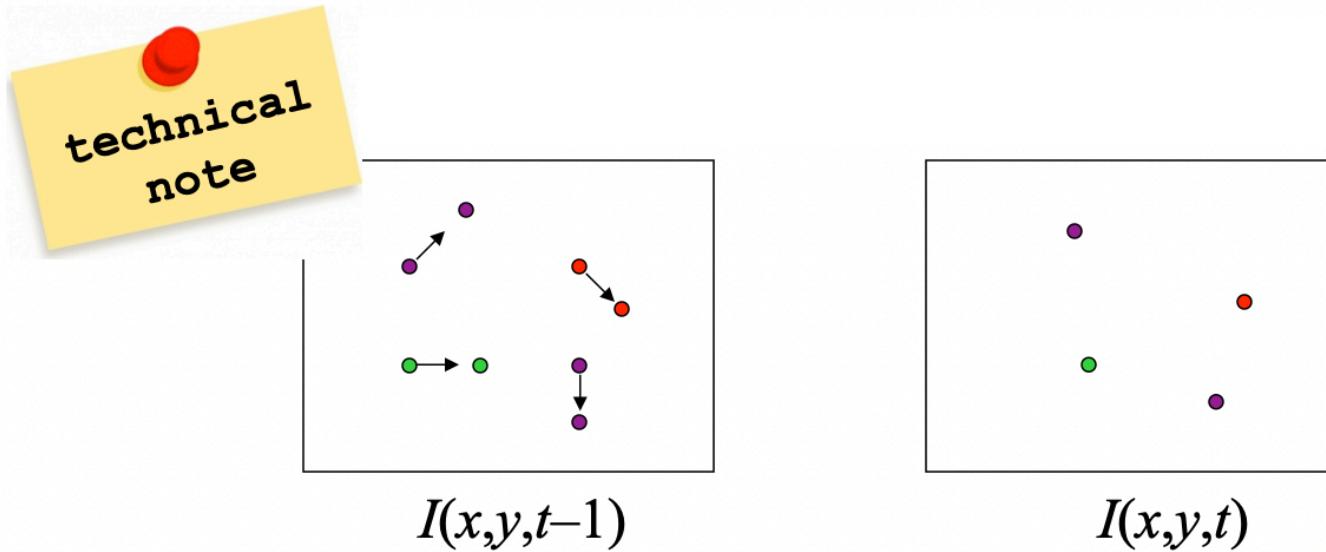
**GOAL:** Recover image motion at each pixel from optical flow

# Optical Flow



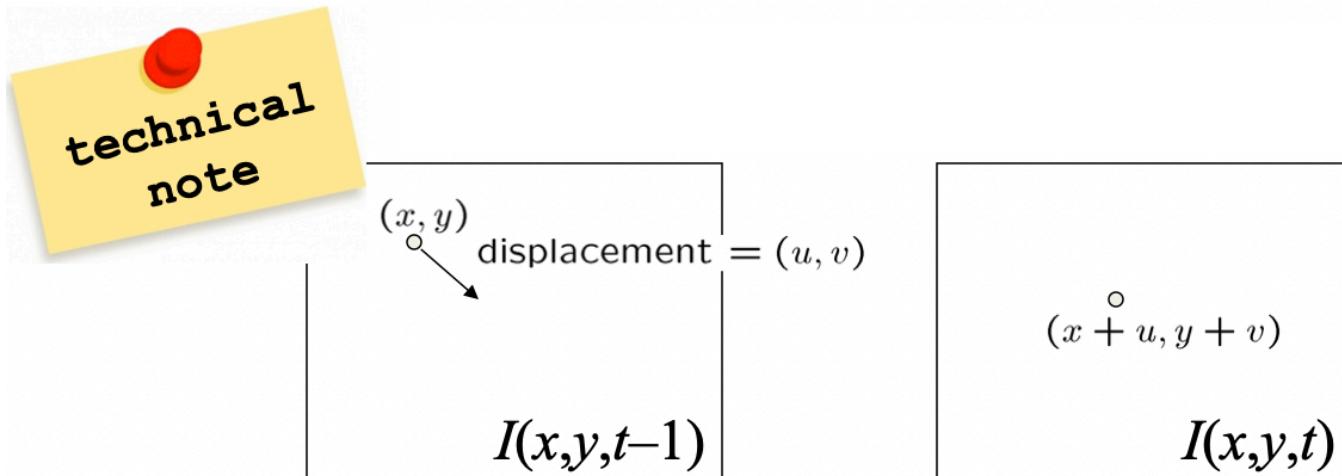
Picture courtesy of Selim Temizer - Learning and Intelligent Systems (LIS) Group, MIT

# Estimating Optical Flow



- Given two subsequent frames, estimate the apparent motion field  $u(x,y), v(x,y)$  between them
- Key assumptions
  - **Brightness constancy:** projection of the same point looks the same in every frame
  - **Small motion:** points do not move very far
  - **Spatial coherence:** points move like their neighbors

# The Brightness Constancy Constraint



- Brightness Constancy Equation:

$$I(x, y, t - 1) = I(x + u(x, y), y + v(x, y), t)$$

Linearizing the right side using Taylor expansion:

$$I(x + u, y + v, t) \approx I(x, y, t - 1) + I_x \cdot u(x, y) + I_y \cdot v(x, y) + I_t$$

Image derivative along x

$$I(x + u, y + v, t) - I(x, y, t - 1) = I_x \cdot u(x, y) + I_y \cdot v(x, y) + I_t$$

$$\text{Hence, } I_x \cdot u + I_y \cdot v + I_t \approx 0 \rightarrow \nabla I \cdot [u \ v]^T + I_t = 0$$

# The Brightness Constancy Constraint

Can we use this equation to recover image motion ( $u, v$ ) at each pixel?

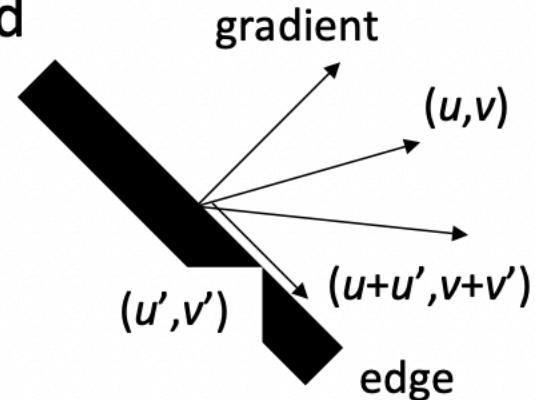
$$\nabla I \cdot [u \ v]^T + I_t = 0$$

- How many equations and unknowns per pixel?
  - One equation (this is a scalar equation!), two unknowns ( $u, v$ )

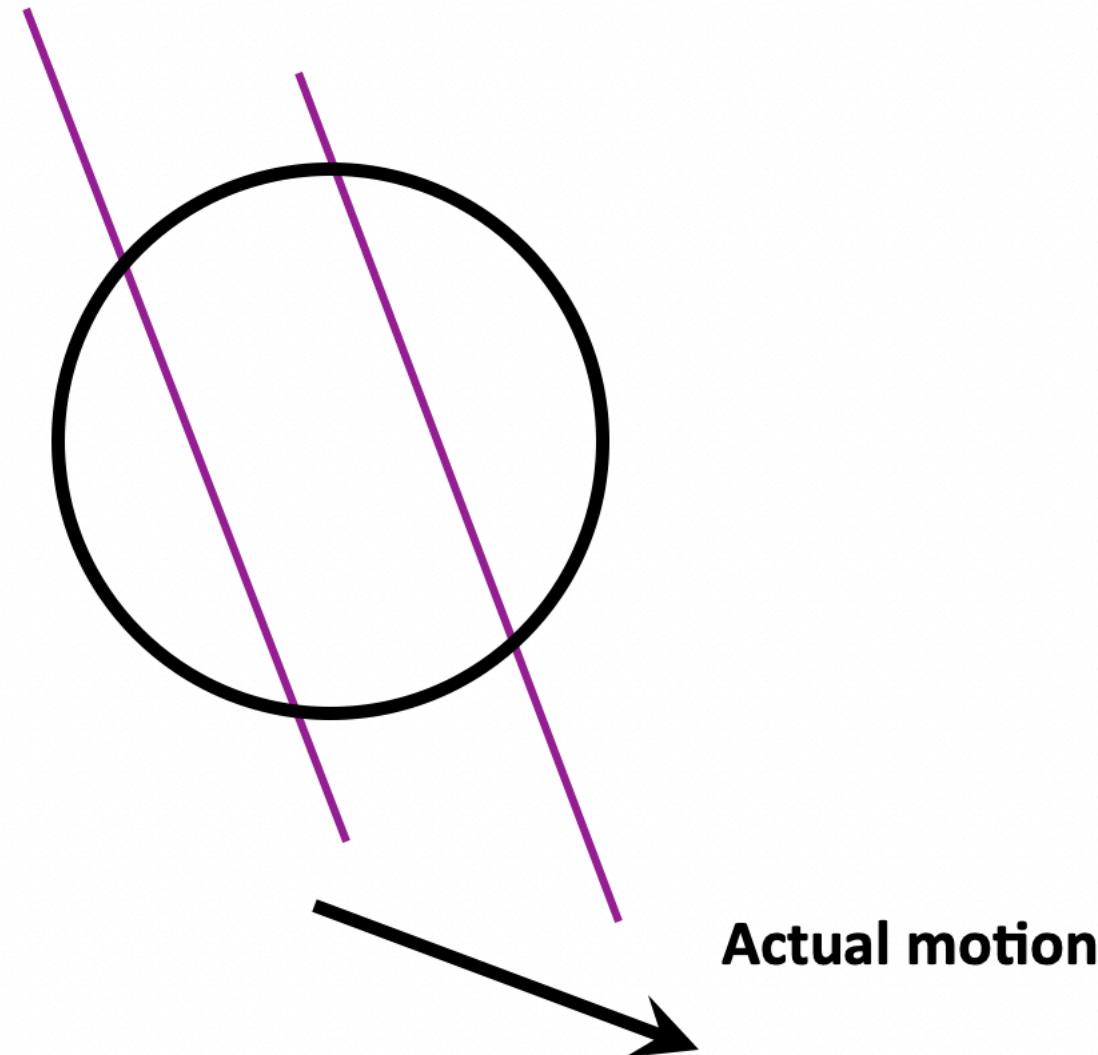
The component of the flow perpendicular to the gradient (i.e., parallel to the edge) cannot be measured

If  $(u, v)$  satisfies the equation,  
so does  $(u+u', v+v')$  if

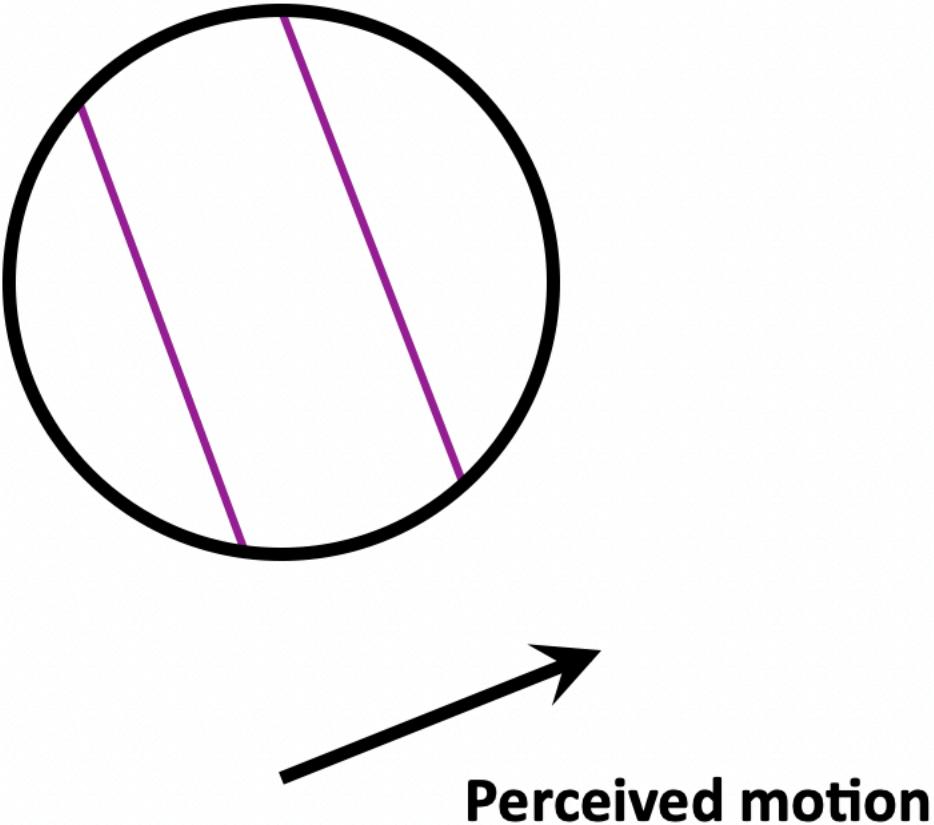
$$\nabla I \cdot [u' \ v']^T = 0$$



# The Aperture Problem



# The Aperture Problem



# Barberpole Illusion



- [https://en.wikipedia.org/wiki/Barberpole\\_illusion](https://en.wikipedia.org/wiki/Barberpole_illusion)

# Solving the Ambiguity

- How to get more equations for a pixel?
- **Spatial coherence constraint:**
- Assume the pixel's neighbors have the same  $(u, v)$ 
  - If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

B. Lucas and T. Kanade. An iterative image registration technique with an application to stereo vision. In Proceedings of the International Joint Conference on Artificial Intelligence, pp. 674- 679, 1981.

# Lucas-Kanade Flow

- Overconstrained linear system:

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

$A \quad d = b$   
 $25 \times 2 \quad 2 \times 1 \quad 25 \times 1$

# Lucas-Kanade Flow

- Overconstrained linear system:

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

$A \quad d = b$   
 $25 \times 2 \quad 2 \times 1 \quad 25 \times 1$

Least squares solution for  $d$  given by  $(A^T A)^{-1} A^T b$

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A \qquad \qquad \qquad A^T b$

The summations are over all pixels in the  $K \times K$  window

# Condition for Solvability

– Optimal  $(u, v)$  satisfies Lucas-Kanade equation

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A$                                      $A^T b$

## When is This Solvable?

- $A^T A$  should be invertible
- $A^T A$  should not be too small due to noise
  - eigenvalues  $\lambda_1$  and  $\lambda_2$  of  $A^T A$  should not be too small
- $A^T A$  should be well-conditioned
  - $\lambda_1 / \lambda_2$  should not be too large ( $\lambda_1$  = larger eigenvalue)

Does this remind anything to you?

# Condition for Solvability

$M = A^T A$  is the *second moment matrix* !

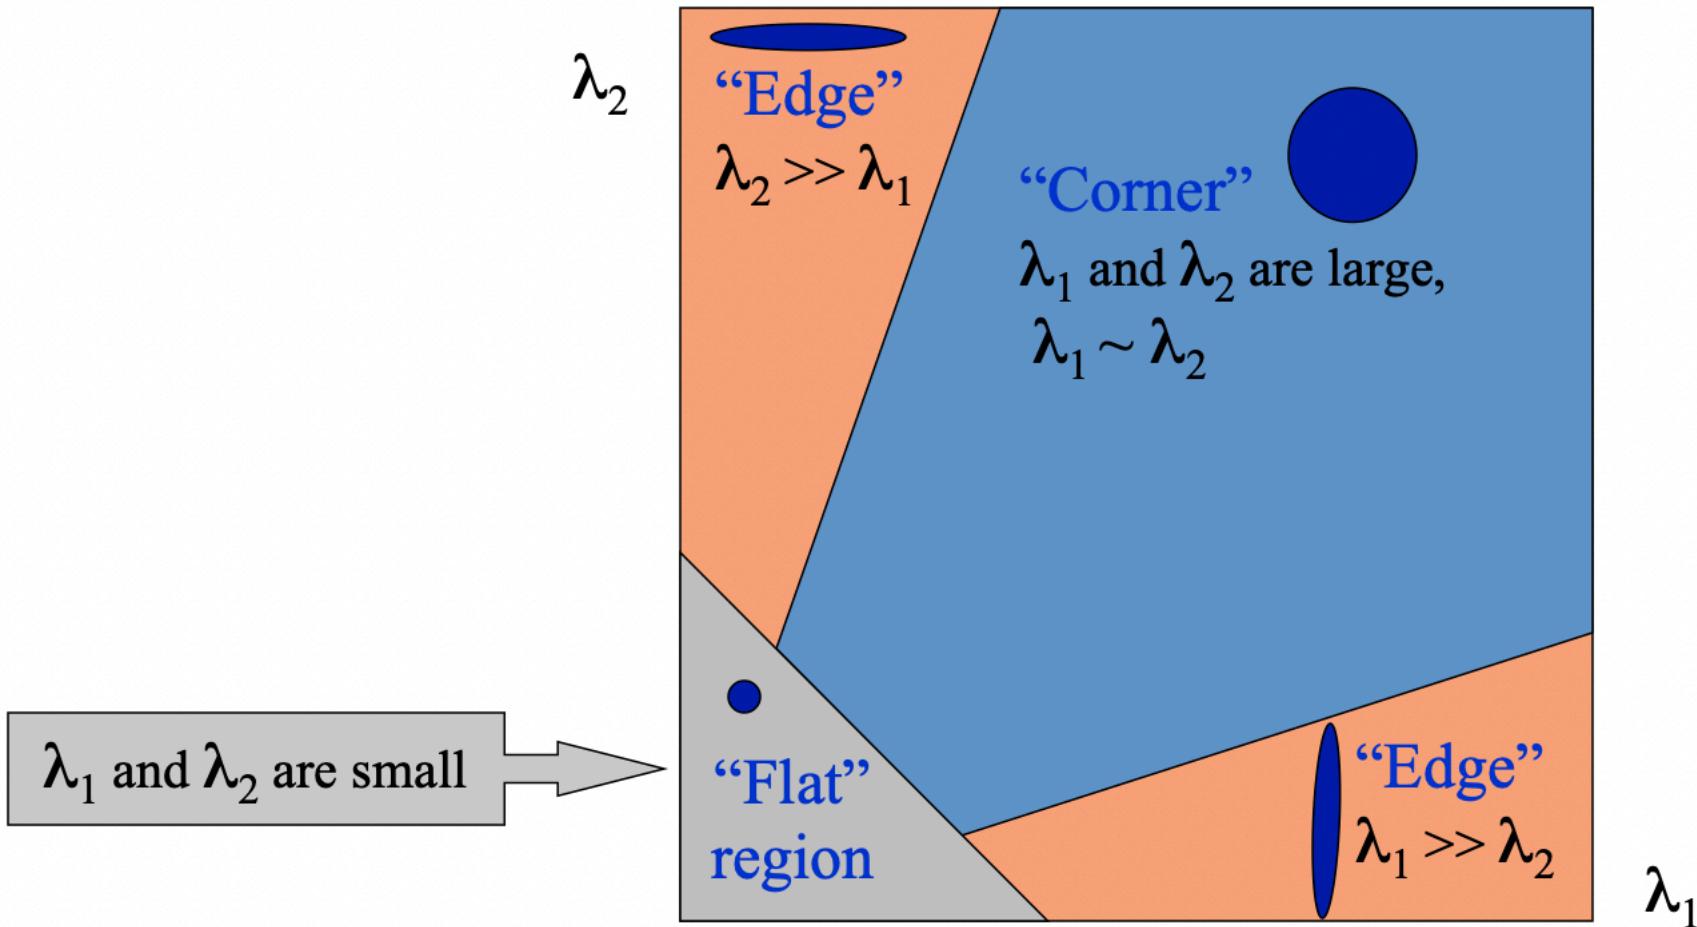
(Harris corner detector...)

$$A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

- Eigenvectors and eigenvalues of  $A^T A$  relate to edge direction and magnitude
  - The eigenvector associated with the larger eigenvalue points in the direction of fastest intensity change
  - The other eigenvector is orthogonal to it

# Interpreting the Eigenvalues

Classification of image points using eigenvalues of the second moment matrix:



# Edge



$$\sum \nabla I (\nabla I)^T$$

- gradients very large or very small
- large  $\lambda_1$ , small  $\lambda_2$

# Low-Texture Region



$$\sum \nabla I (\nabla I)^T$$

- gradients have small magnitude
- small  $\lambda_1$ , small  $\lambda_2$

# High Texture Region



$$\sum \nabla I (\nabla I)^T$$

- gradients are different, large magnitudes
- large  $\lambda_1$ , large  $\lambda_2$

# Measuring Motion

Image at frame  $t$

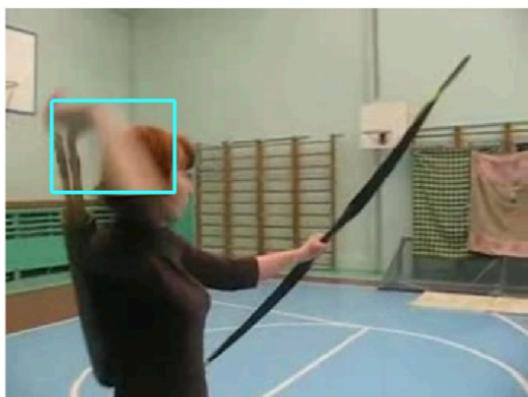
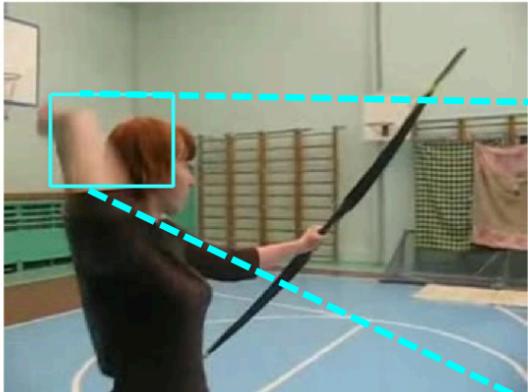
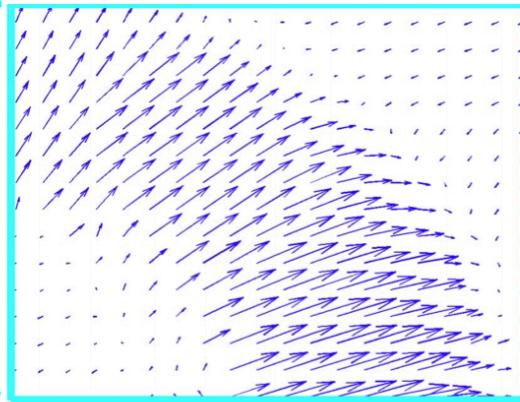


Image at frame  $t+1$

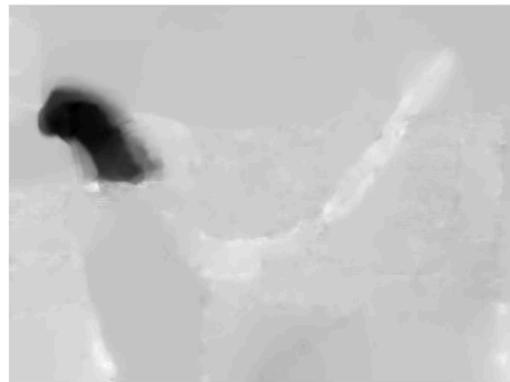
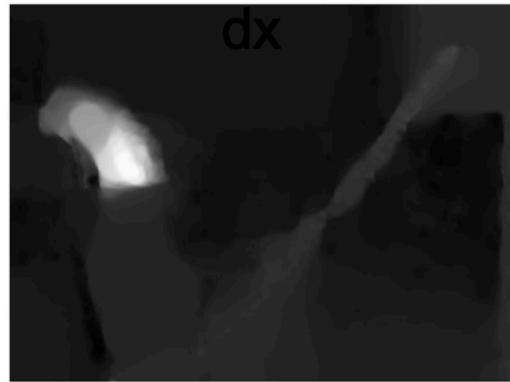
Optical flow gives a displacement field  $F$  between images  $I_t$  and  $I_{t+1}$



Tells where each pixel will move in the next frame:  
 $F(x, y) = (dx, dy)$   
 $I_{t+1}(x+dx, y+dy) = I_t(x, y)$

Optical Flow highlights  
**local motion**

Horizontal flow



Vertical Flow  $dy$

Slide credit: Justin Johnson

# FlowNet: Learning Optical Flow with Convolutional Networks

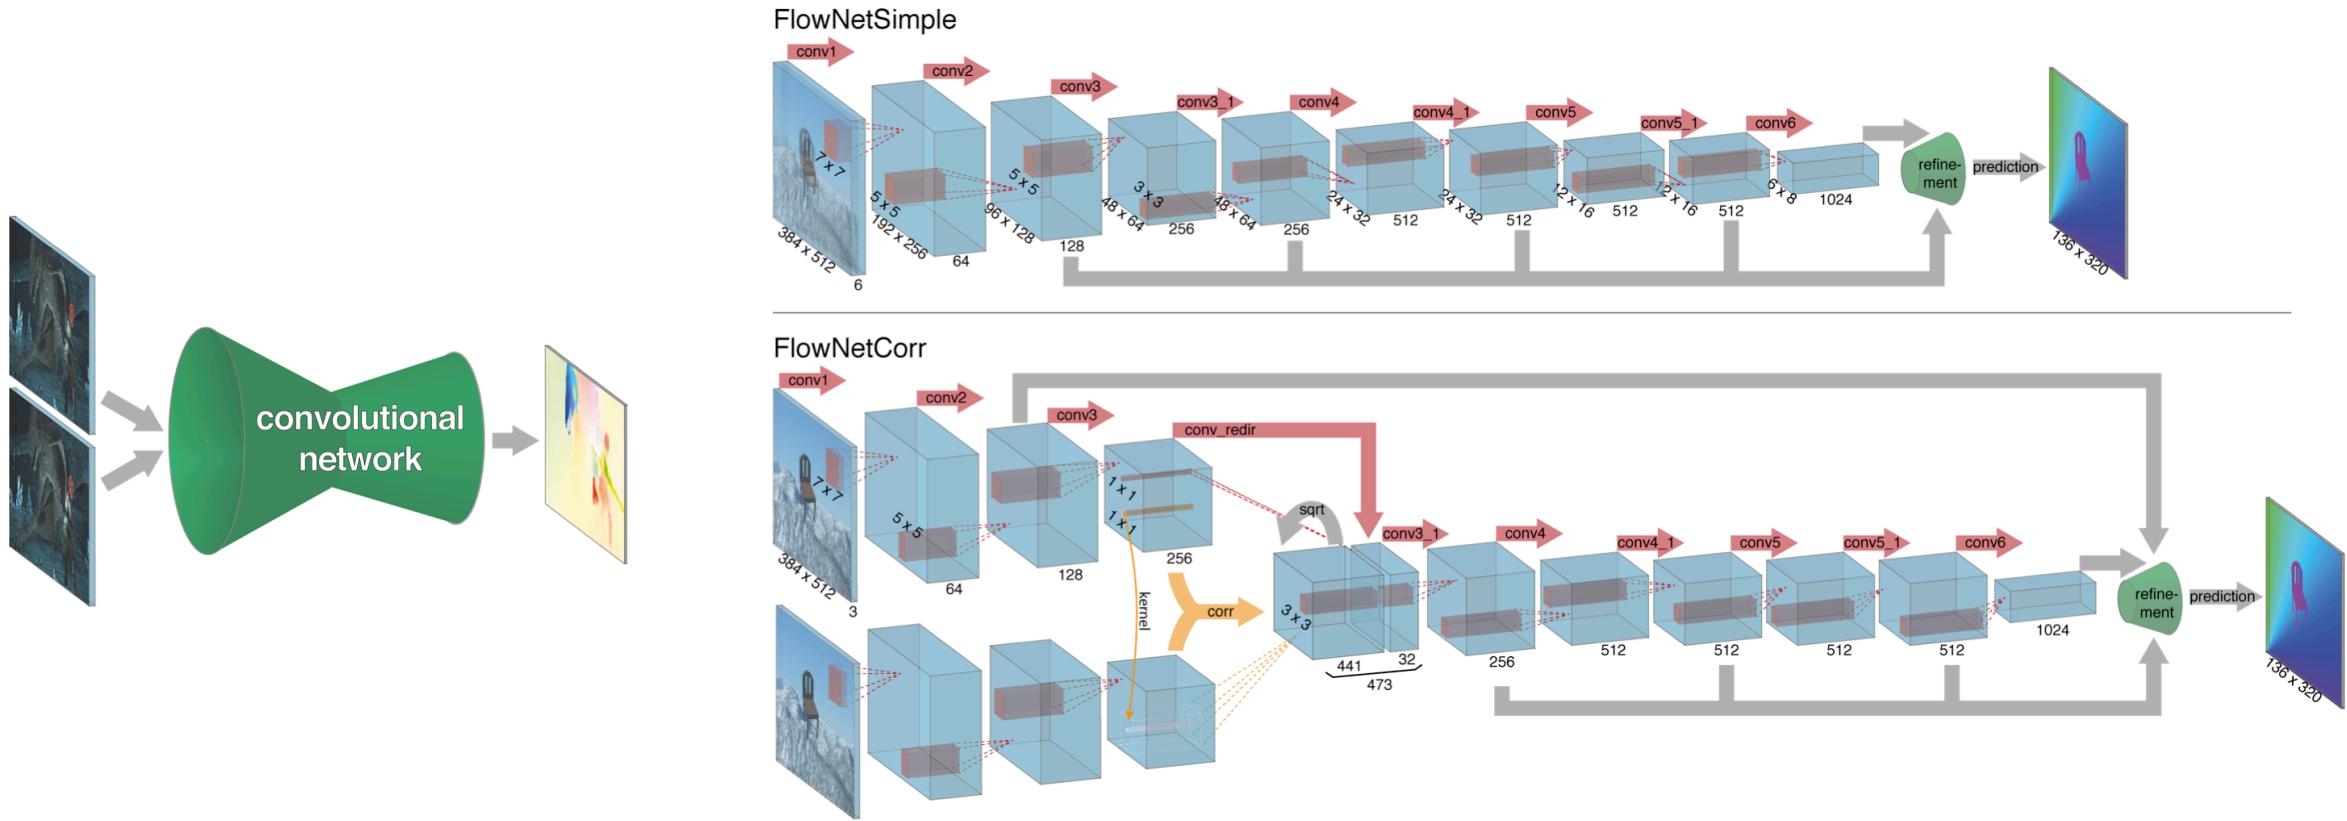


Figure 2. The two network architectures: FlowNetSimple (top) and FlowNetCorr (bottom).



# Introduction to Computer Vision

Next week:  
Lecture 11 Temporal Analysis II