



# Operating Systems (A) (Honor Track)

---

## Lecture 11: Windows Virtual Memory

Yao Guo (郭耀)

Peking University

Fall 2024

Acknowledgements: Prof. Xiangqun Chen & Tao Wang at PKU and Prof. Yuanyuan Zhou at UCSD

# Buzz Words

---

**Protected Mode**

**Virtual Address  
Descriptor (VAD)**

**Working set**

**Self-Map**



# This Lecture

---

## Windows Virtual Memory

Intel x86 Virtual Memory

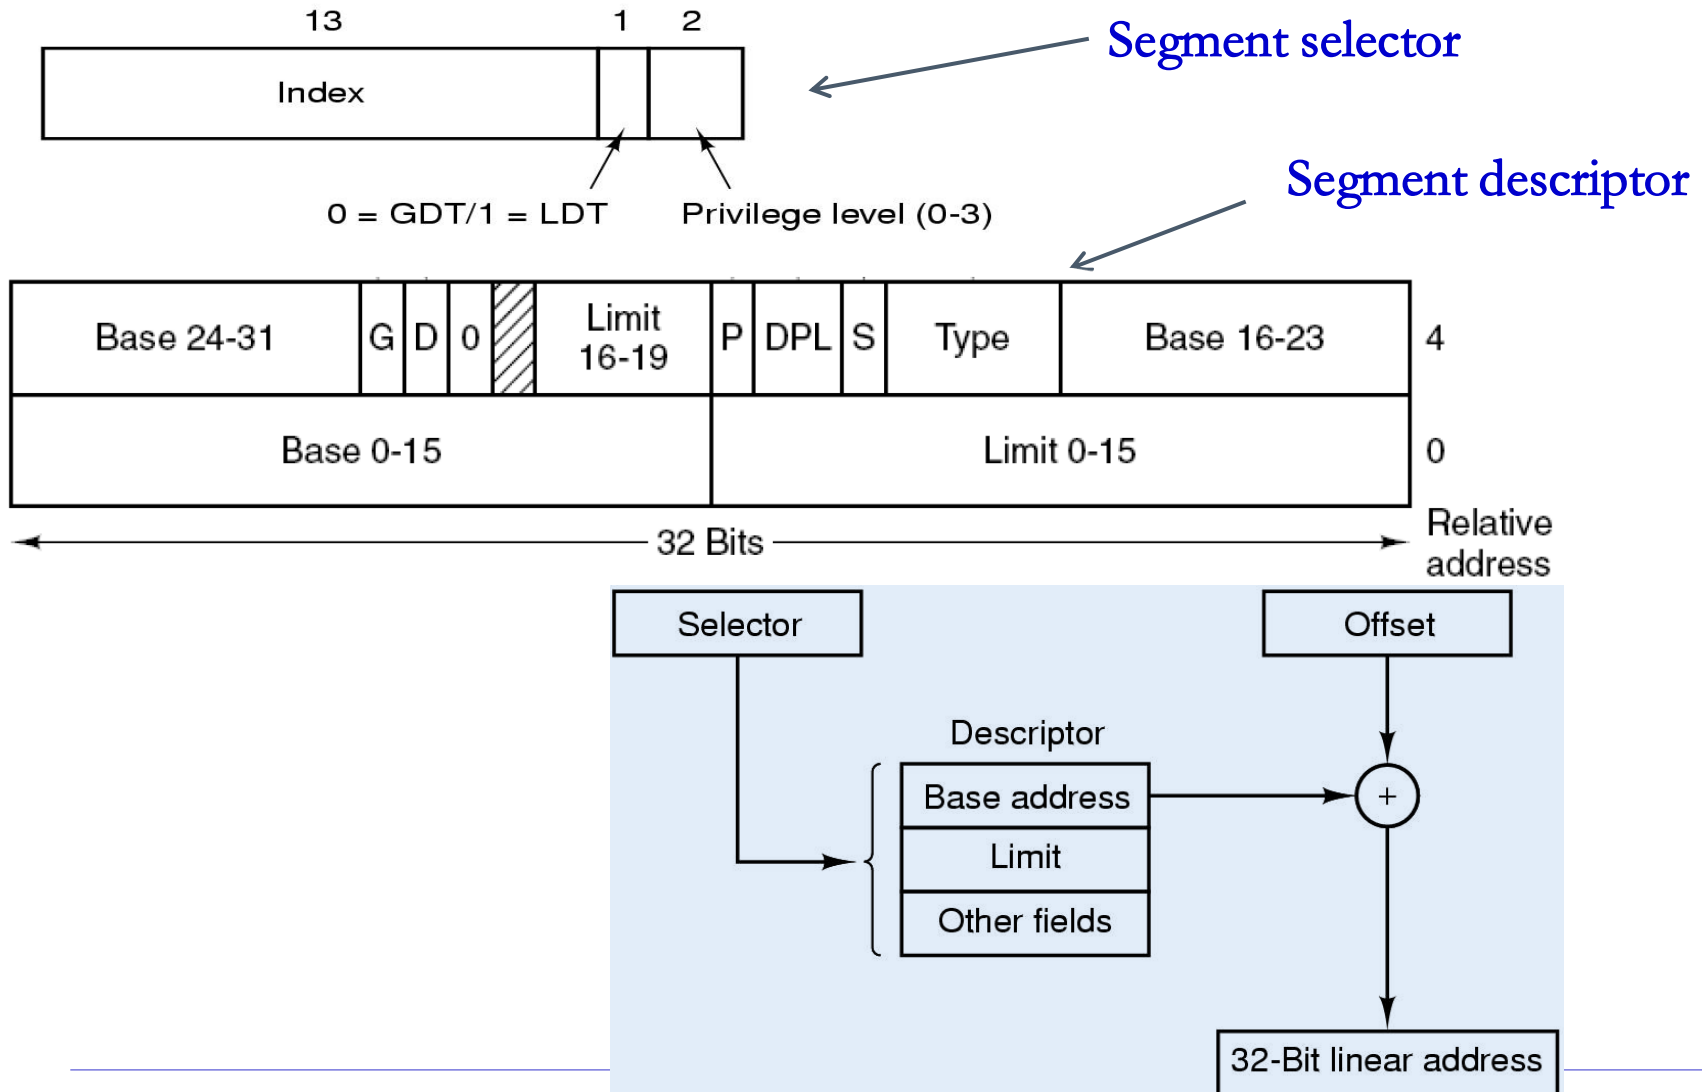
Windows MM

# Address Mode for x86


---

- Real mode (8086)
  - Using physical address
  - For DOS
- Protected mode (80286)
  - Protected virtual address mode
  - Supports virtual memory, paging and safe multi-tasking
  - Supports both 16 and 32 bits
- Intel 8086 virtual mode (80386)
  - Hardware virtualization
  - Run real mode in the protected mode


# Protected Mode



# Logical to Physical Address



(a)



# i386 Page Table Entries

## PDE (Page Directory Entry)

PFN	Avail	G	P S	0	A	PCD	PWT	U/S	R / W	P
-----	-------	---	--------	---	---	-----	-----	-----	-------------	---

## PTE (Page Table Entry)

PFN	Avail	G	0	D	A	PCD	PWT	U/S	R / W	P
-----	-------	---	---	---	---	-----	-----	-----	-------------	---

63	62	52	51	12	11	9	8	7	6	5	4	3	2	1	0
N X	AVL	Physical page number			AVL	G	P A T	D	A	P C D	P W T	U / S	R / W	P	

NX No eXecute

AVL AVaiLable to the OS

G Global page

PAT Page Attribute Table

D Dirty (modified)

A Accessed (referenced)

PCD Page Cache Disable

PWT Page Write-Through

U/S User/Supervisor

R/W Read/Write access

P Present (valid)

**Intel x86 or  
AMD x64 PTEs**



# This Lecture

---

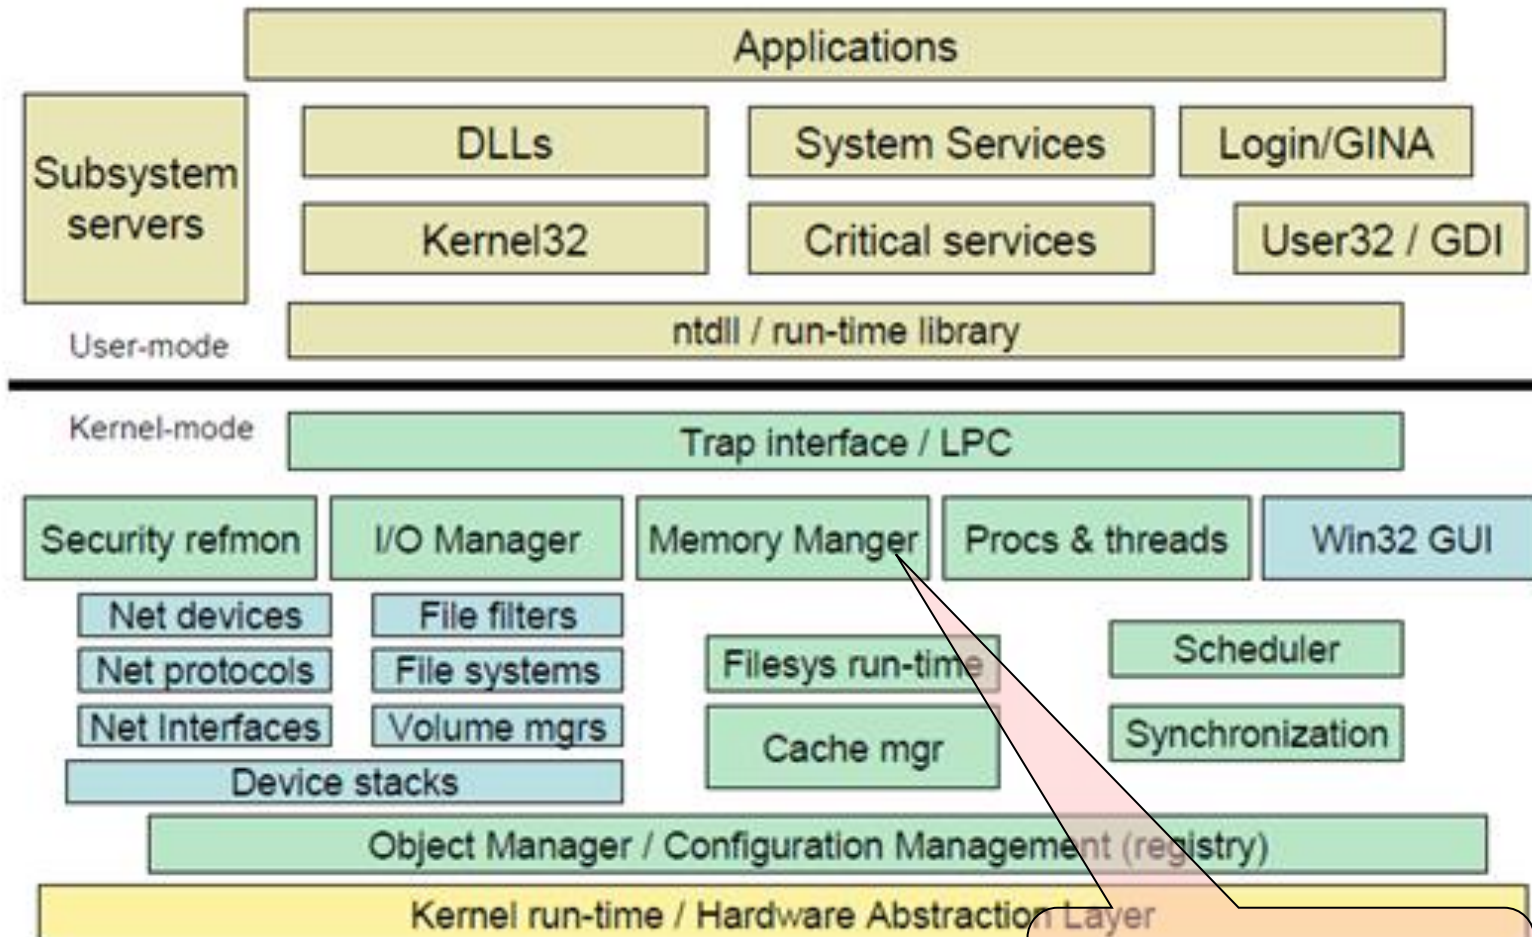
## Windows MM

Intel x86 Virtual Memory

Windows MM



# The Windows Architecture



**MM:**  
**Ntoskrnl.exe**

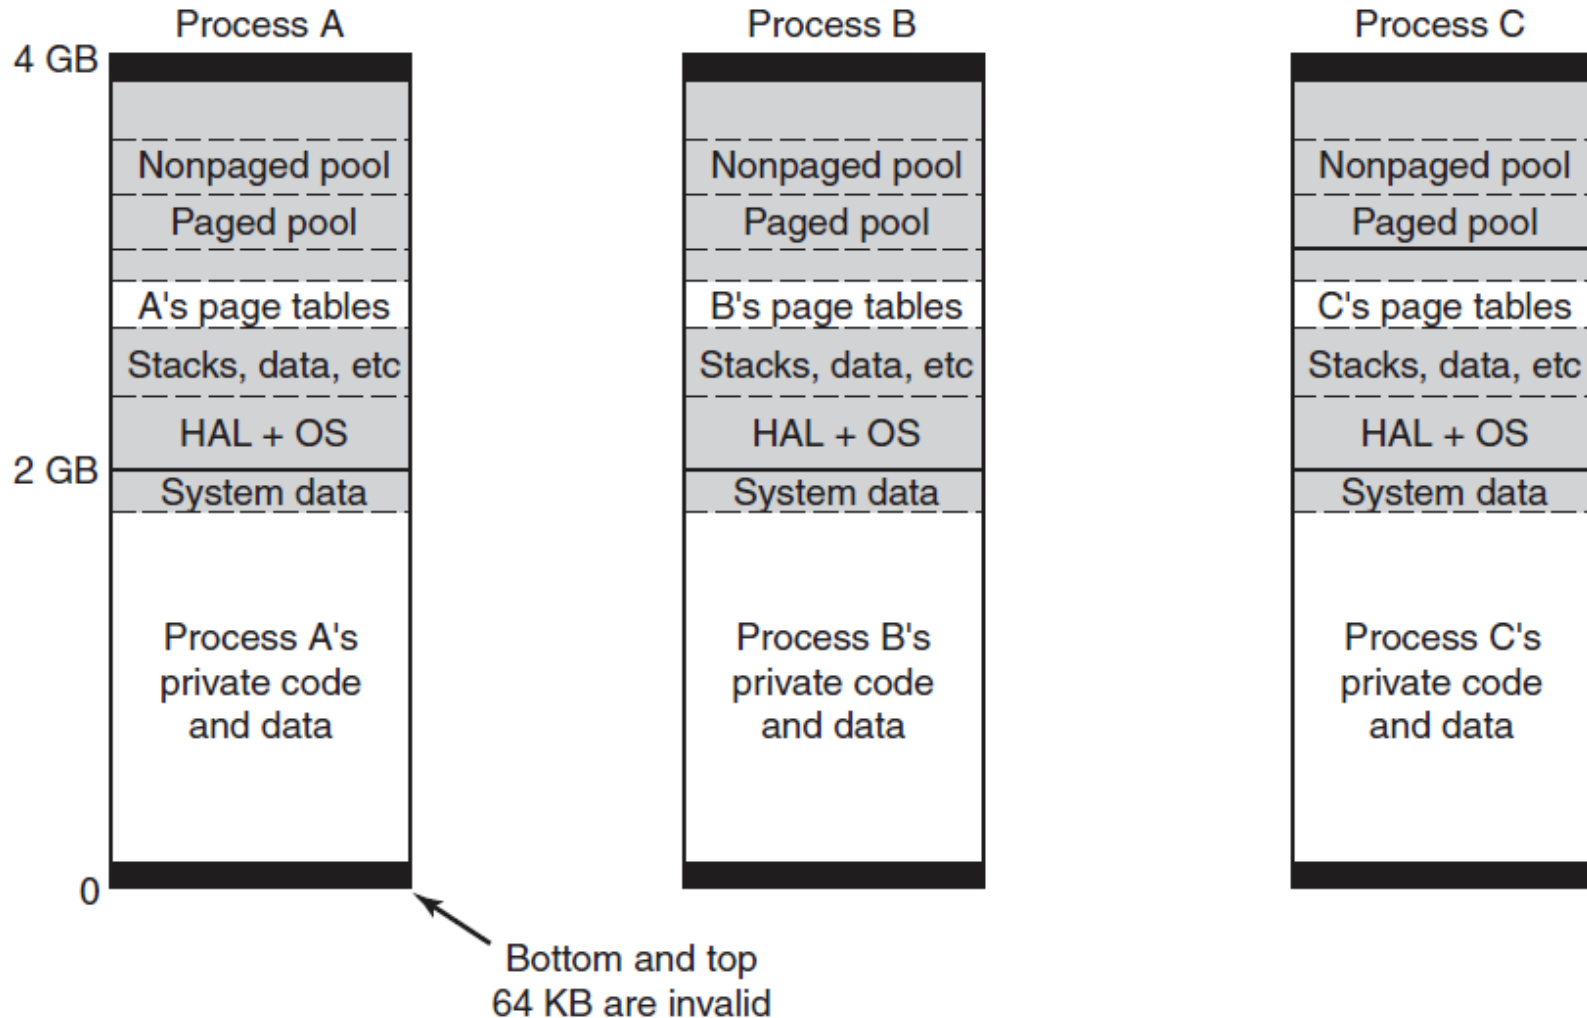


# Windows Virtual Memory

---

- x86
  - Virtual addresses: 32 bits
  - Address space for each process: 4GB
    - Kernel: 2GB
    - User : 2GB
- x86-64
  - Large enough
- Demand paging
  - Page size: 4KB
    - (2MB pages are also used)

# Virtual Address Space Layout





# Virtual Address Allocation

---

- Each page of virtual addresses can be in one of three states:
  - **Invalid**: not currently mapped to a memory section object and a reference to it causes a page fault
  - **Committed**: code or data is mapped onto a virtual page
  - **Reserved**: invalid but has the property that those virtual addresses will never be allocated by the memory manager for another purpose
    - function as **guard pages** to keep the stack from growing too far and overwriting other process data.

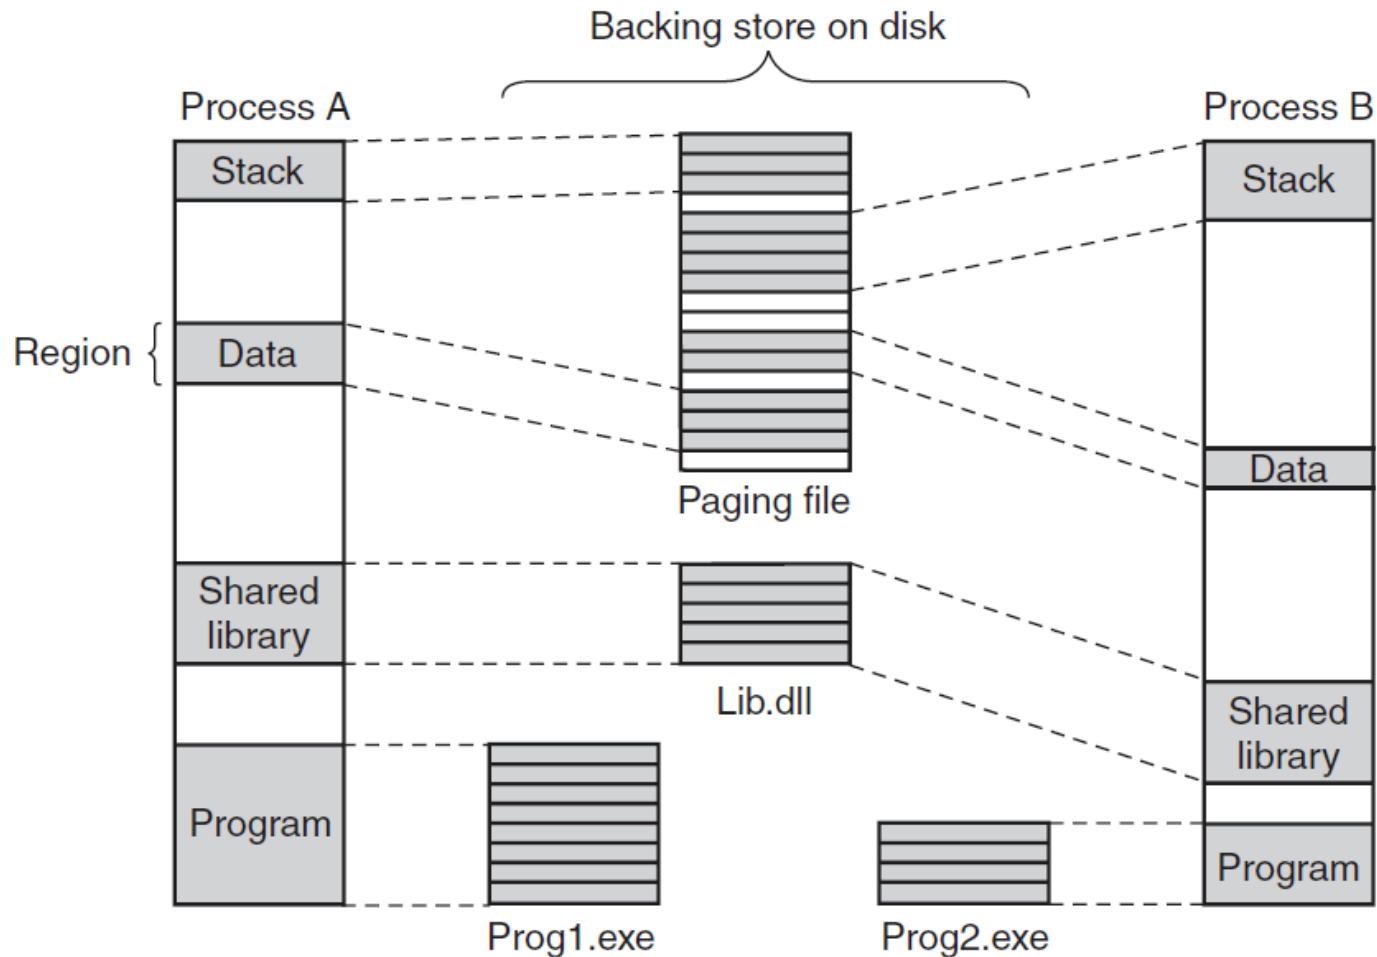


# Windows MM System Calls

---

Win32 API function	Description
VirtualAlloc	Reserve or commit a region
VirtualFree	Release or decommit a region
VirtualProtect	Change the read/write/execute protection on a region
VirtualQuery	Inquire about the status of a region
VirtualLock	Make a region memory resident (i.e., disable paging for it)
VirtualUnlock	Make a region pageable in the usual way
CreateFileMapping	Create a file-mapping object and (optionally) assign it a name
MapViewOfFile	Map (part of) a file into the address space
UnmapViewOfFile	Remove a mapped file from the address space
OpenFileMapping	Open a previously created file-mapping object

# Windows MM Implementation



**Figure 11-30.** Mapped regions with their shadow pages on disk. The *lib.dll* file is mapped into two address spaces at the same time.

# VAD (Virtual Address Descriptor)

---

- MM creates a VAD (Virtual Address Descriptor) for each process
  - Listing the range of addresses mapped, the section representing the backing store file and offset where it is mapped, and the permissions
  - An address space is completely defined by the list of its VADs
  - VADs are organized into a balanced tree

# Page Fault Handling

---

- Types of pages faults:
  - The page referenced is not committed.
  - Access to a page has been attempted in violation of the permissions.
  - A shared copy-on-write page was about to be modified.
  - The stack needs to grow.
  - The page referenced is committed but not currently mapped in.
  
- How does each of these types occur?




# Hard Faults vs. Soft Faults

---

- **Hard fault:** needs reading from disk
- **Soft fault:** the memory manager can satisfy a page fault by finding the needed page in memory rather than reading it in from disk
- **When might a soft fault occur?**
  - A shared page has already been mapped into another process
  - Only a new zero page is needed
  - The needed page was trimmed from the process' working set but is being requested again before it has had a chance to be reused
  - Pages have been compressed to effectively increase the size of physical memory.

# Self-Map Entries

- The **self-map entries** are used to map the physical pages of the page tables and page directory into kernel virtual addresses (shown for 32-bit PTEs).



Self-map: PD[0xc0300000>>22] is PD (page-directory)

Virtual address (a): (PTE \*) (0xc0300c00) points to PD[0x300] which is the self-map page directory entry

Virtual address (b): (PTE \*) (0xc0390c84) points to PTE for virtual address 0xe4321000

# Self-Map Entries

---

**MiGetPdeAddress():**

Given a virtual address  $va$ , compute its PDE

$((\text{PMMPTE})((((\text{ULONG})(va)) \gg 22) \ll 2) + \text{PDE\_BASE}))$

**MiGetPteAddress():**

Given a virtual address  $va$ , compute its PTE

$((\text{PMMPTE})((((\text{ULONG})(va)) \gg 12) \ll 2) + \text{PTE\_BASE}))$



# Page Replacement Algorithm

---

## □ Working set

- Each process' working set is described by two parameters: the minimum size and the maximum size.
  - Default minimum: 20–50 pages
  - Default maximum: 45–345 pages

## □ The working set manager runs every second

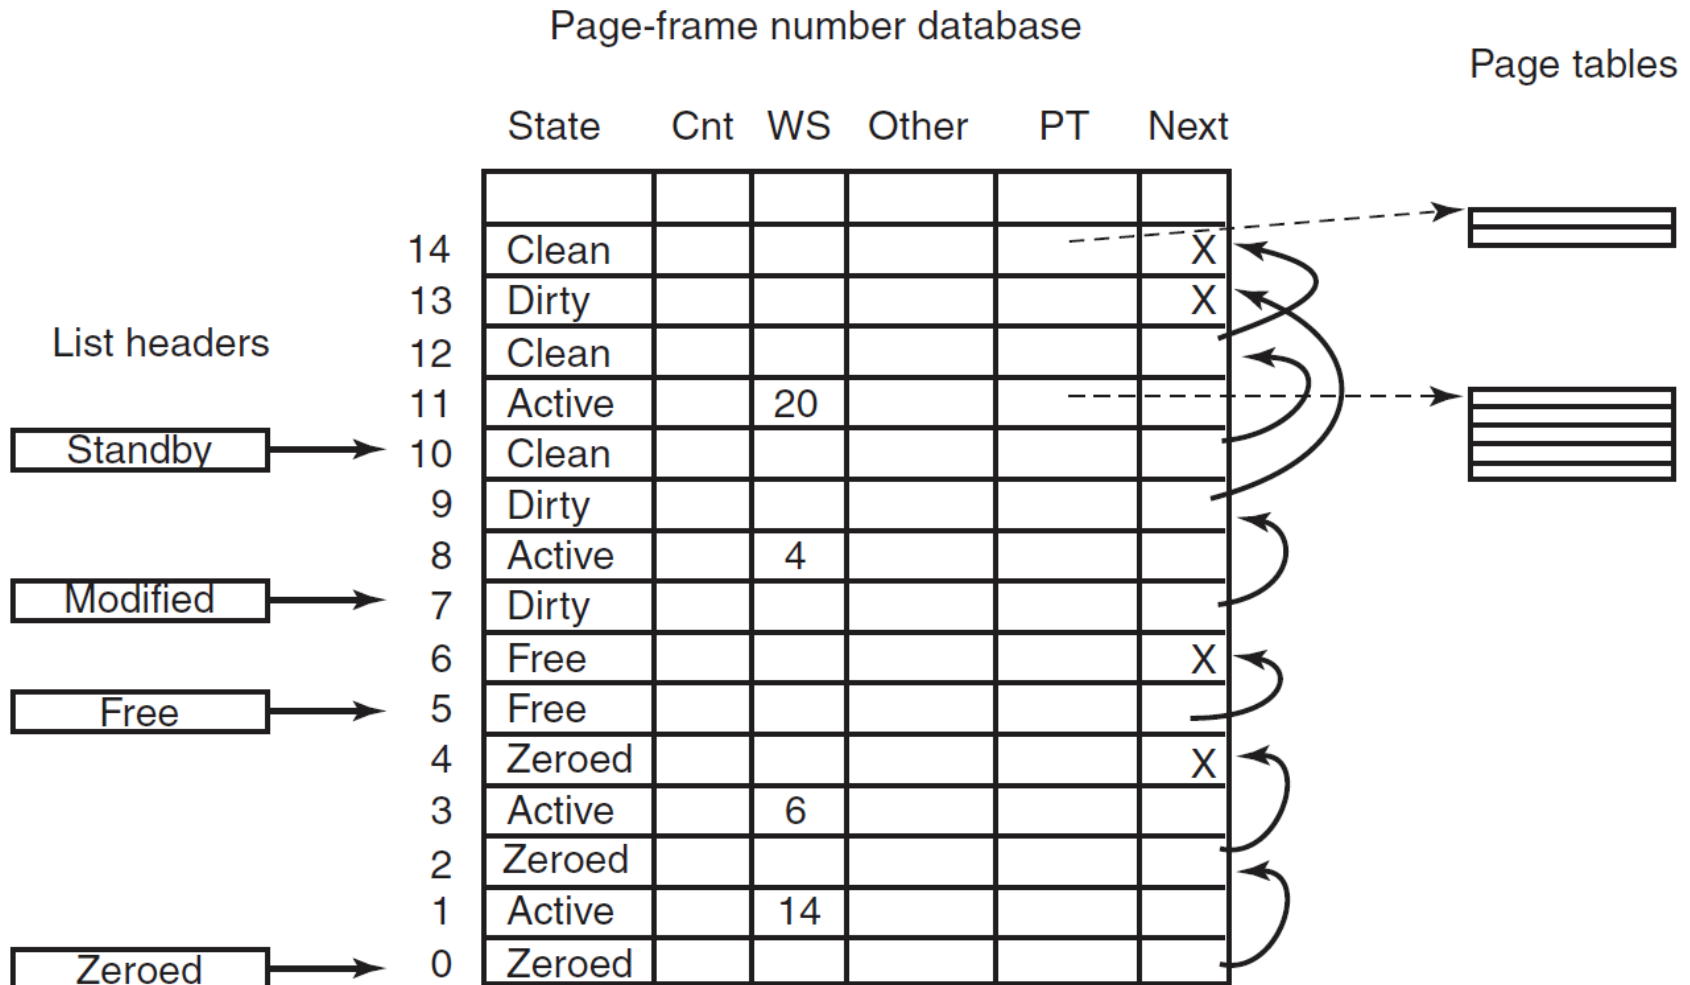
- The working-set manager throttles the amount of work it does to keep from overloading the system.

# Physical Memory Management

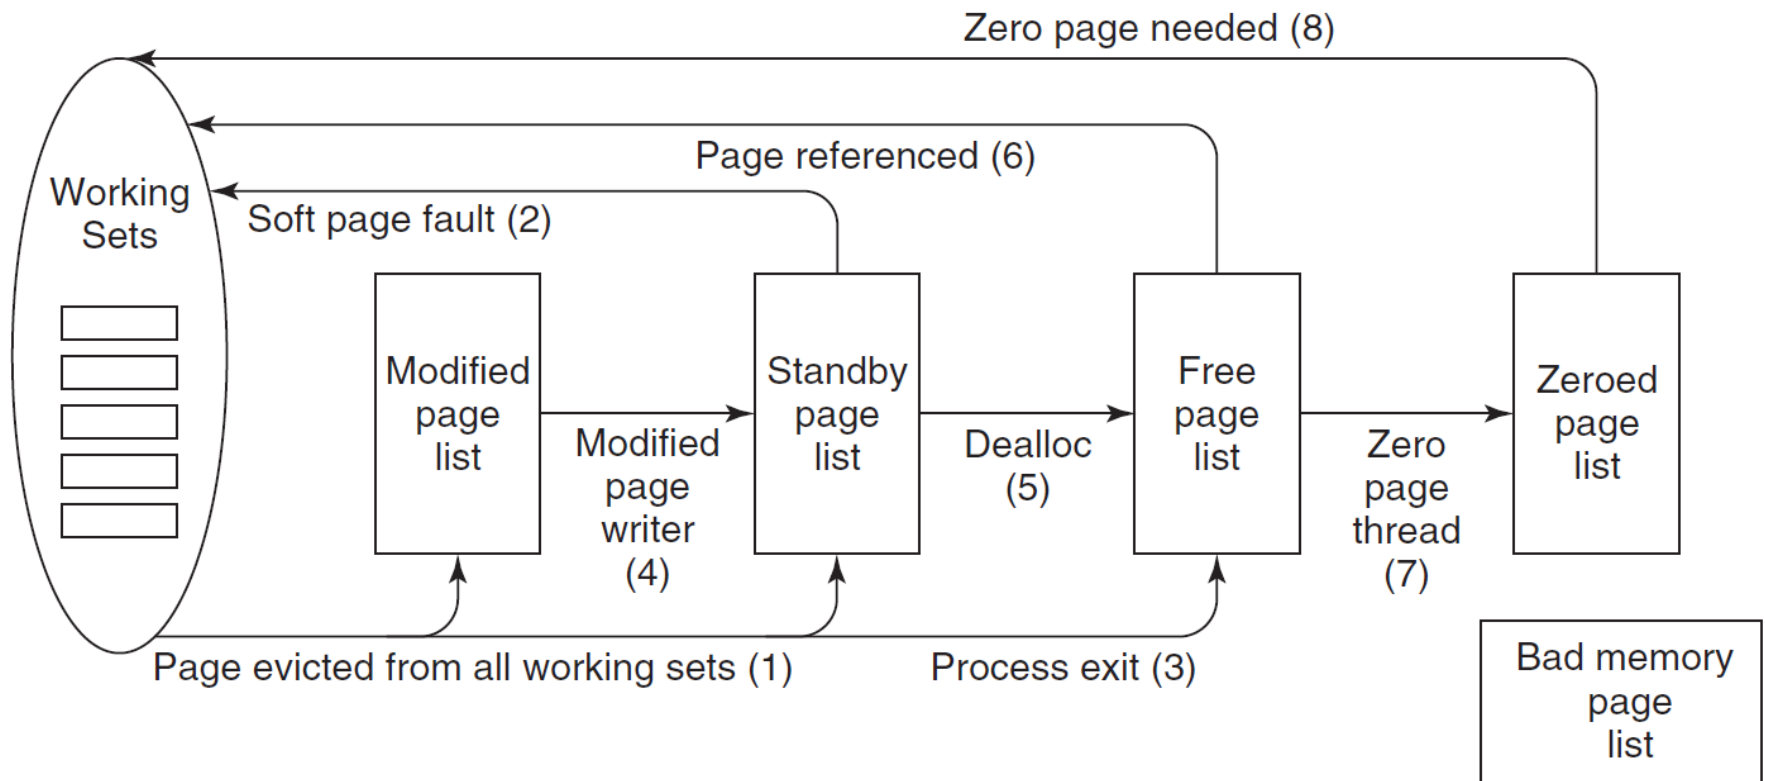
---

- PFN (Page Frame Number) Database
  - All pages in the system either are referenced by a valid page-table entry or are on one of these five lists
  
- Different page frame lists
  - The free list
  - The standby list
  - The modified list
  - The zeroed list (free & zeroed)
  - The bad memory page list (frames with hardware errors)

# Physical Memory Management



# Transitions between Page Lists



**Figure 11-34.** The various page lists and the transitions between them.

# Summary

---

- Windows Memory Management
  - Virtual Address Descriptor
  - Self-Map Entries
  - Various Kinds of Page Lists
  
- Next Lecture: Scheduling